

Manual - Running the Account Enrichment Steps

Living document. We extend this manual as new steps are added. *Last updated: 16 June 2026*

What this pipeline does

Once a year we export the full CRM account list and run it through three small scripts that automatically add two kinds of information:

- **Segment** - what kind of organisation each account is (Local government, Healthcare, IT & software, ...). See step 1b.
- **Population** (`cx_population`) - the number of inhabitants, for the government bodies among the accounts. See step 2.

The result is one Excel file (`final_enriched.xlsx`) that can be loaded back into Dynamics 365. The rest of this manual is the step-by-step run guide; this opening section explains **which accounts get a population number and where that number comes from**.

Which accounts get a population, and from where

A population is only filled for organisations that actually have inhabitants (a municipality does; an IT company does not). The classifier (step 1) labels every account with a type; step 2 then uses the source below per type. If a lookup finds nothing, the account keeps its existing CRM value (it is never blanked).

Account type (example)	Country	Population comes from
Gemeente / municipality (Gemeente Amsterdam)	NL	CBS "Gebieden in Nederland" - the official statistics office, current-year figure. Merged/abolished gemeenten fall back to their last-known figure from Wikidata.
Gemeente / municipality (Gemeente Gavere)	BE	Wikidata (Belgian municipality list). Merged ones (e.g. Zwiindrecht, 2025) fall back to their last-known figure.
Gemeinde / municipality (Stadt Lübeck)	DE	Wikidata, looked up via the official municipality key (AGS). Same-name towns are told apart by the account's postal code.
OCMW (OCMW Gavere)	BE	Population of the gemeente it belongs to (same source as BE gemeente).
AGB - autonomous municipal company (AGB Mortsel)	BE	Population of the gemeente it belongs to.
Landratsamt / Kreisverwaltung (Landratsamt Kusel)	DE	Population of the Landkreis it belongs to.
Provincie / province (Provincie Utrecht)	NL	Wikidata (NL province list).
Provincie / province (Provincie Antwerpen)	BE	Wikidata (BE province list).
Landkreis / district (Landkreis Bautzen)	DE	Wikidata, looked up via the official district key.
Politiezone / police zone (Politiezone Antwerpen)	BE	Sum of the populations of the gemeenten the zone covers (173 zones mapped from the official list).
Veiligheidsregio / safety region (Veiligheidsregio Zeeland)	NL	Sum of its member gemeenten; the membership list comes from CBS (all 25 regions).

Account type (example)	Country	Population comes from
Omgevingsdienst / environmental agency	NL	Sum of its member gemeenten (mapping partially filled).
Waterschap / water authority (Waterschap Aa en Maas)	NL	Fixed table of the 21 water authorities, taken from their own websites.
Stadsdeel / city district (Stadsdeel Centrum)	NL	Fixed table of the 8 Amsterdam districts, from Wikipedia.
Verbandsgemeinde (Verbandsgemeinde Daun)	DE	Fixed table from the German Wikipedia (official Landesamt figures).
Land / country government (Caribbean) (Bestuur Aruba)	AW/CW/SX/BQ	Wikidata country total (Aruba, Curaçao, Sint Maarten, Caribisch Nederland).

Accounts that deliberately get NO population (they have no inhabitants): commercial companies, ministries / Rijksoverheid, Belgian FOD, Stadtwerke, intercommunales, CAW, Zweckverband. These keep an empty population field.

Types not yet automated (they keep their existing CRM value until we build a mapping for them): inter-municipal cooperations (`samenwerking` , `belastingsamenwerking`), GGD health regions, `stadsregio` , BE `hulpverleningszone` , German `Amt` and `Verwaltungsgemeinschaft` . These are relatively few accounts; see the maintenance section for how to add them.

*Every output row also has a `bron` column naming the exact source used, and a `proces` column explaining the decision - so any single value can be traced back. The technical view of these sources (which code table, how to refresh) is in *"How the population values are sourced"* further down.*

Getting the input from the data warehouse (step 0)

- 1 Log into Microsoft SQL Server Management Studio (DWH).
- 2 Connect to the DWH database.
- 3 Go to Views -> `INT.Account` .
- 4 Run:

```
SELECT *
FROM [INT].[Account]
WHERE statecode = 0
```

The `statecode = 0` filter keeps only **active** accounts; disabled accounts are excluded from the pipeline.

- 5 Right-click the results grid, choose **"Save Results As..."** and save as `full_input.csv` .
- 6 Copy `full_input.csv` to the project folder on the Linux server (`~/get-populations/` , see one-time setup below).

The scripts read this CSV directly (UTF-8 with BOM, `NULL` text values are handled). If disabled records accidentally end up in the export, step 1 logs a warning.

The pipeline at a glance

The scripts run as a chain. **Order matters**: each step's output is the next step's input, so all added columns accumulate into one final file.

```
full_input.csv (DWH export, see step 0)
|
| python step1_classify.py --input full_input.csv --output step1_classified.xlsx
v
step1_classified.xlsx          (+ type/country detection columns)
|
| python step1b_segment.py --input step1_classified.xlsx --output step1b_segmented.xlsx --web-search
v
step1b_segmented.xlsx          (+ Segment, Segment (detailed) columns)
```

```
|
| python step2_enrich.py --input step1b_segmented.xlsx --output final_enriched.xlsx \
|     --user-agent "your-org/1.0 (you@company.com)"
|
v
final_enriched.xlsx          (+ population columns) <- THE final file
```

Step 1 is free and instant. Step 1b uses the Claude API (costs money, uses the cache). Step 2 uses Wikidata (free; the first run downloads reference data, which takes 5-10 minutes, after that it is cached in `reference/`).

Every run of every step also writes a timestamped log file to `logs/` (e.g. `logs/step1b_segment_2026-06-11_143052.log`), so a failed or unattended run can be diagnosed afterwards. The `logs/` folder stays out of git.

Which steps need the paid Claude API?

Step	Needs
step 0 (DWH export)	Nothing (internal DWH access)
step 1 (classify)	Nothing - runs locally, free
step 1b (segment)	Anthropic API key + prepaid credits
step 2 (population)	Nothing - Wikidata is free (just put a real contact email in <code>--user-agent</code>)

Important: a Claude.ai subscription (Pro or Max) does **not** cover API usage. The API is billed separately via prepaid credits, even if you already pay for a Claude subscription. Claude Code / claude.ai usage and API usage are two different wallets.

Getting an API key and credits (one-time)

- 1 Go to console.anthropic.com and sign in (or create an account).
- 2 Under **Billing / Plans & billing**, buy prepaid credits with a credit card. \$50 comfortably covers the first full run plus re-runs; yearly refreshes after that cost a few dollars at most.
- 3 Under **API Keys**, create a new key and copy the `sk-ant-...` value (it is shown only once).
- 4 Store it on the server:

```
echo 'export ANTHROPIC_API_KEY="sk-ant-..." >> ~/.bashrc
```

0. One-time setup (first time on a machine)

This assumes a fresh Linux box with `git` and `python3` installed and nothing of this project yet. We install everything under the home directory (`~/get-populations`); no root access is needed.

```
# 1. get the code from GitHub
cd ~
git clone https://github.com/Golden-VS/get-populations.git
cd get-populations

# 2. create a Python environment and install all dependencies
python3 -m venv venv
source venv/bin/activate
pip install pandas openpyxl anthropic requests rapidfuzz
```

Set the Anthropic API key (from `console.anthropic.com`, where the credits live - only step 1b needs it):

```
export ANTHROPIC_API_KEY="sk-ant-..."
```

To avoid retyping it every session, store it once in your shell profile:

```
echo 'export ANTHROPIC_API_KEY="sk-ant-..." >> ~/.bashrc
```

Before every working session, go to the folder and activate the environment first:

```
cd ~/get-populations
source venv/bin/activate
```

1. Classification (step 1)

```
python step1_classify.py --input full_input.csv --output step1_classified.xlsx
```

- Free and fast (a few seconds for ~6,000 records); no API needed.
- Detects per account the entity type (gemeente, waterschap, politiezone, commercial, ...) and the country, based on the account name.
- Check the summary in the terminal (also in the `logs/` file): - the distribution per `detected_type` should look plausible (hundreds of municipalities, no unexpectedly empty types); - the `""Onbekend'` maar wel een `cx_businessstype""` list and the `*conflicts*` table show records worth a quick manual glance.
- Output `step1_classified.xlsx` is the input for step 1b.

2. Segmentation test run (step 1b, 100 accounts)

Always do this first after any change to the script or the input file.

```
python step1b_segment.py \
  --input step1_classified.xlsx \
  --output step1b_test.xlsx \
  --test-mode --web-search
```

- Takes a few minutes, costs well under \$1.
- Open `step1b_test.xlsx` and check the columns "Segment", "Segment (detailed)", "Segment (category)" (Governmental / Non-profit / Commercial / Unknown), `segment_confidence` and `segment_bron`.
- Pay extra attention to rows with confidence `low` and segment `Commercial (other)` or `Unknown`.

3. Segmentation full run (step 1b, all ~6,000 accounts)

```
python step1b_segment.py \
  --input step1_classified.xlsx \
  --output step1b_segmented.xlsx \
  --web-search
```

- Takes roughly 1 to 1.5 hours. Progress is logged per batch.
- Estimated cost: \$15 to \$25 including web searches (covered by the \$50 credit purchase).
- Safe to interrupt: results are saved to `segment_cache.csv` after every batch. Re-running the same command resumes where it stopped and does not pay twice for accounts already classified.
- At the end the log prints a count per segment. Sanity-check that distribution (e.g. roughly 1,500 municipalities, no giant "Unknown" bucket).

4. Corrections

If a segment is wrong, do not edit the output Excel by hand (it would be overwritten on the next run). Instead, add a row to a corrections file, for example `overrides_segment.xlsx`, with these columns:

accountid	segment_override	segment_detailed_override	reden
(CRM guid)	IT & software	Software vendor	manually verified

Then re-run with:

```
python step1b_segment.py \
  --input step1_classified.xlsx \
  --output step1b_segmented.xlsx \
  --web-search --overrides overrides_segment.xlsx
```

Overrides always win over the automatic classification and cost nothing.

5. Population enrichment (step 2)

After segmentation, run step 2 on **step1b's output** (not on step1_classified.xlsx, or the Segment columns will be missing from the final file):

```
python step2_enrich.py \
  --input step1b_segmented.xlsx \
  --output final_enriched.xlsx \
  --user-agent "your-org/1.0 (you@company.com)"
```

- Wikidata requires a contact address in `--user-agent`. The whole string is free text in the form `toolname/version (email)` - both parts are yours to choose, nothing is registered anywhere. It only serves to identify the caller so Wikimedia can contact you if a query misbehaves. Concrete example:

```
--user-agent "crm-enrichment/1.0 (jan.jansen@yourcompany.nl)"
```

Pick any sensible tool name, keep `1.0` as version, and use a real work email.

- First run downloads reference data (5-10 min); later runs use the cache in `reference/` (refreshed yearly automatically, or force with `--refresh-cache`).
- Use `--test-mode` here too when running for the first time.
- `final_enriched.xlsx` is the complete file: original columns + type detection + segments + population.

How the population columns work

- `cx_population` is updated *in place* (same column name as in the CRM, so the file can be imported back directly). If no new value is found for a record, the old CRM value is **kept**, never blanked.
- `previous_population` (new column) holds the old CRM value as it was before this run, so you can filter on records where the value changed.
- `population_gewijzigd` (new column): `gewijzigd / gelijk / leeg`. Filter on `gewijzigd` to review exactly which values would change in the CRM; the `run_log` tab also shows the total count.
- Supporting columns explain every value: `bron` (source, incl. URL or Wikidata ID), `proces` (short explanation in Dutch), `peildatum_inwoners` (reference year of the statistic), `data_leeftijd_jaren` (age of that statistic in years - high values flag stale data), `match_score` and `invuldatum`.

Population corrections

Step 2 has its own corrections table (separate from the segment one). Create e.g. `overrides_population.xlsx` with columns:

account_id	population_override	reden
(CRM guid)	154000	figure from annual report 2025

and run with `--overrides overrides_population.xlsx`. Overrides win over every automatic lookup.

How the population values are sourced (technical / maintenance view)

The table at the very top of this manual lists the source **per account type**. This section is the maintenance view: the four underlying **mechanisms** behind those sources and **where each lives in the code**, so you know what to touch when something needs updating.

Mechanism	How it works	Where in the code
1. Wikidata lookup	Downloads reference lists with populations from Wikidata (free, no key) and fuzzy-matches account names. Same-name German municipalities are told apart by postal code; dissolved NL/BE gemeenten fall back to a frozen last-known figure via an exact-name list.	REFERENCE_SOURCES in <code>step2_enrich.py</code> ; cached in <code>reference/*.csv</code>
2. CBS open data	The official "Gebieden in Nederland" table - current-year NL gemeente figures, and the veiligheidsregio membership used by mechanism 3. Auto-discovers the newest yearly edition.	<code>fetch_cbs_gebieden()</code> ; cached in <code>reference/cbs_gebieden.csv</code>
3. Sum of member gemeenten	Entity population = sum of the gemeenten it covers (politiezones, veiligheidsregio's, omgevingsdiensten).	BE_POLITIEZONE_GEMEENTEN , NL_OMGEVINGSDIENST_GEMEENTEN ; veiligheidsregio mapping from CBS at run time (inline table is the fallback)
4. Direct-value tables	Hand-collected numbers with a source URL per entry, where no database has the data (waterschappen, Amsterdam stadsdelen, DE Verbandsgemeinden).	NL_WATERSCHAP_INWONERS , NL_STADSDEEL_INWONERS , DE_VERBANDSGEMEINDE_INWONERS
5. Manual override	A corrections file that wins over everything.	<code>--overrides</code> file (see above)

If none of these produce a value, the old CRM value is kept (never blanked). Every output row's `bron` column says exactly which source produced its value.

When a source fails or looks wrong (yearly maintenance)

Start with the **"Status per bron"** report at the end of the step 2 log: every reference source is listed as OK / leeg (empty) / GEFAALD (failed).

A Wikidata source is empty or failed:

- 1 Delete that source's file from `reference/` and re-run (forces a fresh download; transient server errors are retried automatically).
- 2 Still broken? The Wikidata class may have changed. Test the query on <https://query.wikidata.org/> and check `CHECKPOINT.md` , which documents how every Q-ID/key was found and validated. Expected healthy counts: NL gemeenten ~1,300 (incl. historical), BE gemeenten ~565, NL provinces 12, BE provinces 10, DE gemeinden ~11,400, DE landkreise ~290.

The CBS source fails: the run does not stop - it logs a warning ("cbs_gebieden GEFAALD") and falls back to Wikidata population figures and the built-in veiligheidsregio table. Values are then 1-2 years older but still correct. Remedy: delete `reference/cbs_gebieden.csv` and re-run; if CBS renamed the table series, check `discover_cbs_gebieden_table()` in `step2_enrich.py` .

A count is suddenly way off (e.g. BE gemeenten drops to 50): treat as broken even if the run "succeeds" - the fuzzy matcher will quietly miss records. Same remedy as above.

Direct-value tables age (methods 3): the numbers drift ~1%/year. Refresh by checking each entry's `bron_url` ; for waterschappen there is a built-in sanity check: the 21 values must sum to roughly the NL population (~18 million). A record that says `"niet in ...-tabel"` in its `proces` column means a new entity appeared (e.g. a new Verbandsgemeinde in the CRM): add one line to the relevant table with the figure from the source in the comment above that table.

Mapping tables outdated after a municipal reorganisation (method 2): politiezone mergers happen in BE every year. The `BE_POLITIEZONE_GEMEENTEN` table was extracted from the NL Wikipedia page "Lijst van politiezones in België"; the extraction recipe is documented in `CHECKPOINT.md` and can be re-run.

General health checks after any run:

- `proces` column: filter on "geen match" / "niet in" to see what was missed.
- `proces` column: filter on "LET OP" to find name collisions that could not be resolved (e.g. a German municipality name that exists in several states while the account has no postal code) - verify those by hand.
- `data_leeftijd_jaren` column: high values (5+) flag stale source data.
- `previous_population` vs `cx_population` : large jumps deserve a look.

6. Next year's run (yearly refresh)

- 1 Make a fresh DWH export following **step 0** at the top of this manual and place `full_input.csv` in the project folder.
- 2 Run the full chain from the pipeline diagram: step 1 -> step 1b (with `--overrides ...` if you have a corrections file) -> step 2.
- 3 Thanks to `segment_cache.csv` , only **new accounts and accounts whose name changed** are sent to the AI. Everything else comes from the cache. Expected cost: a few dollars at most, usually cents.
- 4 Step 2's `reference/` cache is older than 365 days by then, so it re-downloads fresh Wikidata population data automatically.

Do **not** delete `segment_cache.csv` ; it is the memory of all previous classifications. If you ever want to force a complete re-classification (for example after a major change to the segment list), run once with `--refresh-segments` and expect full-run costs again.

Useful options (step 1b)

Option	What it does
<code>--test-mode</code>	Only the first 100 accounts
<code>--web-search</code>	Second pass with web lookup for weak classifications (recommended). Only searches records that were never websearched before, so re-runs stay cheap
<code>--offline</code>	No API calls: mapping table + cache only
<code>--model <id></code>	Use a different Claude model (default <code>claude-opus-4-7</code>)
<code>--overrides <file></code>	Apply a manual corrections file
<code>--refresh-segments</code>	Ignore the cache, classify everything again

Future steps (to be added to this manual)

- Importing the result columns into Dynamics 365